


Losses and Activations

University of Victoria - PHYS 555

Classification

Binary classification recap

- $0 < \sigma < 1$
- Useful to squash layer output to represent binary probability
→ Bernoulli output distribution
- Expensive to compute
- Saturates at low and high input values → small slopes → low gradient signal → needs a log in the loss function to cancel the effect of the exp



$$\sigma(x) = \frac{1}{1 + \exp^{-x}}$$

Softmax function


$$\text{softmax}(\mathbf{x}) = \frac{1}{\sum_{i=1}^n e^{x_i}} \cdot \begin{bmatrix} e^{x_1} \\ e^{x_2} \\ \vdots \\ e^{x_n} \end{bmatrix}$$

- vector of values in $[0,1]$
- behaves like a probability: denominator sums to 1
- is nicely differentiable
- the pre-activated function are called "logits"

Classification and softmax regression

- Multinoulli output distribution \rightarrow multi-class output
- Produces a distribution over classes
- Predicted class is the one with the largest probability

Multi-Class Classification with NN and SoftMax Function



Multi-class loss

- For a classifier with C classes, the network output is: $(f_1(\mathbf{x}; \mathbf{w}), \dots, f_C(\mathbf{x}; \mathbf{w}))$
- We can interpret each output to be a probability, i.e. for NN: $f_j(\mathbf{x}; \mathbf{w}) = \mathbf{x}^T \mathbf{w}_j$

$$P(y = j \mid \mathbf{x}) = \frac{e^{\mathbf{x}^T \mathbf{w}_j}}{\sum_{k=1}^K e^{\mathbf{x}^T \mathbf{w}_k}}$$

- A natural candidate for the loss is using negative log-likelihood:

$$nll = -\frac{1}{N} \sum_n \log P(y = j | \mathbf{x}_n)$$

Cross-entropy for multiclass

- Recall cross-entropy definition between 2 distributions \mathbf{p} and \mathbf{q} :

$$H(p, q) = - \sum_x p(x) \log q(x)$$

- We can then write: $-\log P(Y = j | \mathbf{x}_n) = H(\delta(j), P(Y | \mathbf{x}_n))$

By minimizing our loss, we minimize the average cross-entropy between a deterministic posterior class probability and the estimated class probability

PyTorch Classification Losses Summary

[BCEWithLogitsLoss](#)(\hat{y} , y) = [BCELoss](#)(`sigmoid(\hat{y})`, y)


[CrossEntropyLoss](#)(\hat{y} , y) = [NLLLoss](#)(`log(softmax(\hat{y}))`, y)

Loss function	Usage	Example <i>Using probabilities</i>	Example <i>Using logits</i>
<i>BCELoss or BCEWithLogitsLoss</i>	Binary classification	<i>BCELoss</i> y_{true} : 1 y_{pred} : 0.8	<i>BCEWithLogitsLoss</i> y_{true} : 1 y_{pred} : 0.8
<i>NLLLoss or CrossEntropyLoss</i>	Multiclass classification	<i>NLLLoss</i> y_{true} : 2 y_{pred} : 0.30 0.15 0.55	<i>CrossEntropyLoss</i> y_{true} : 2 y_{pred} : 1.5 0.8 2.1

Regression

Activation in regression problems

- For continuous targets, we do not need a nonlinear function == Identity
- However beware of high dynamic range of vastly different values in your target variables.
 - Use a transformation if this is the case



MSE and MAE

- L2 loss or Mean Square Error:
[MSELoss](#)

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - f(\mathbf{x}_i, \theta))^2$$

- L1 loss or Mean Absolute Error
[L1Loss](#)

$$\text{MAE} = \frac{1}{N} \sum_{i=1}^N |y_i - f(\mathbf{x}_i, \theta)|$$

More robust to outliers
slightly less stable

Quiz

1. How do you transform a one label supervised regression problem into a classification one?
2. How do you transform a binary classification into a regression problem?