

# Sequences in Deep Learning

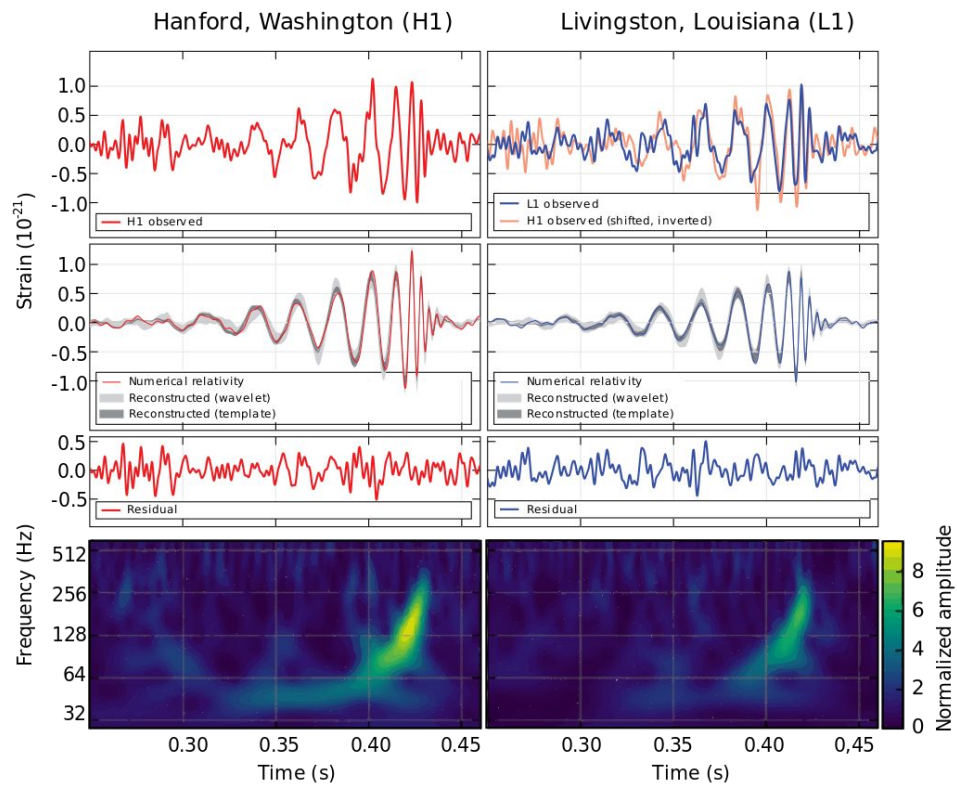
University of Victoria - PHYS-555

# Stock Market

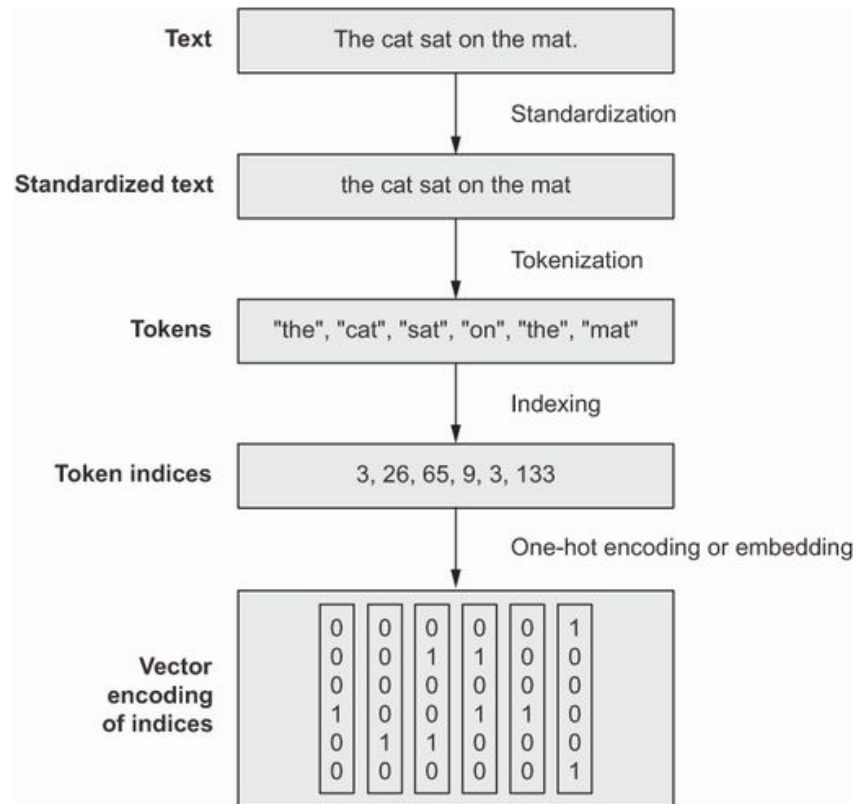
DJIA History 2017-2020



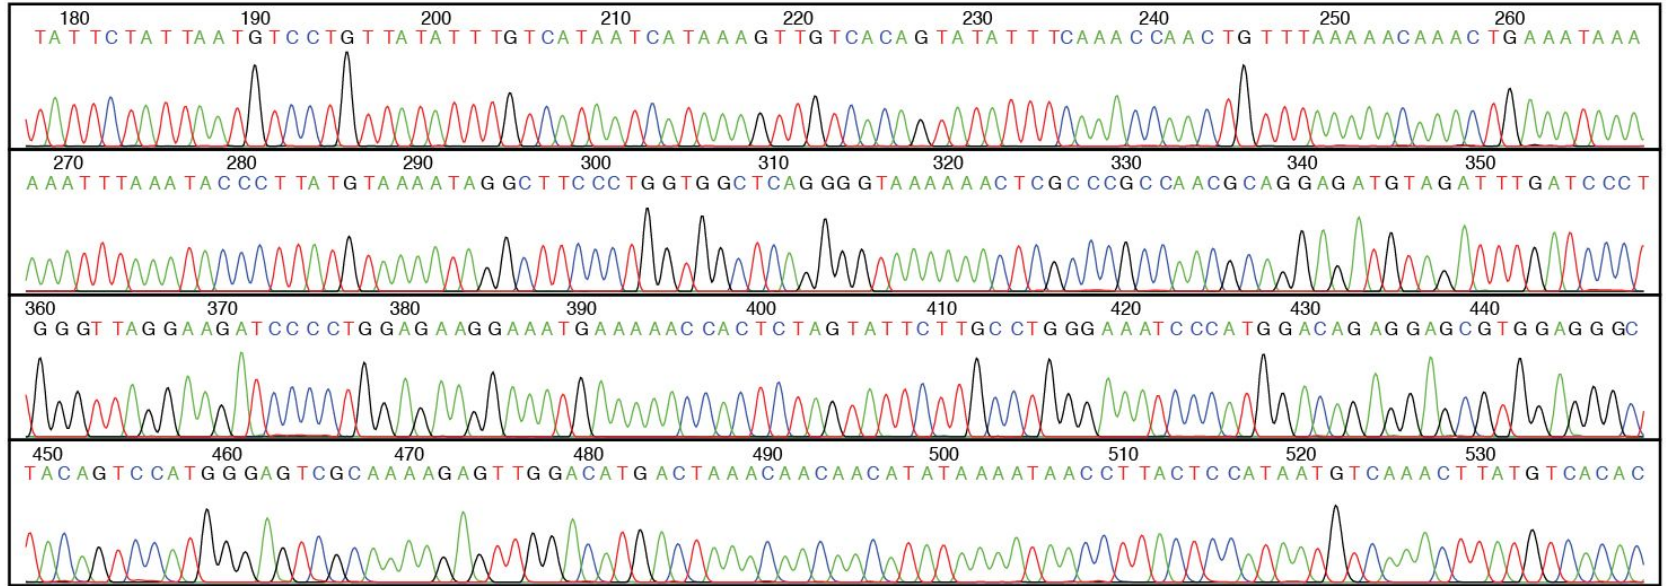
# Gravitational Waves



# Text

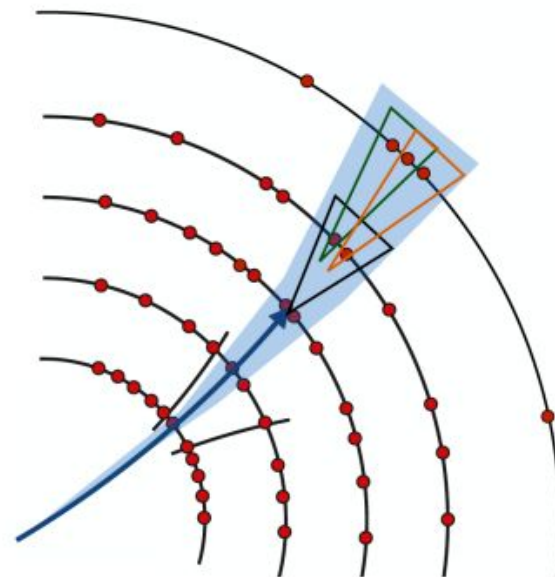
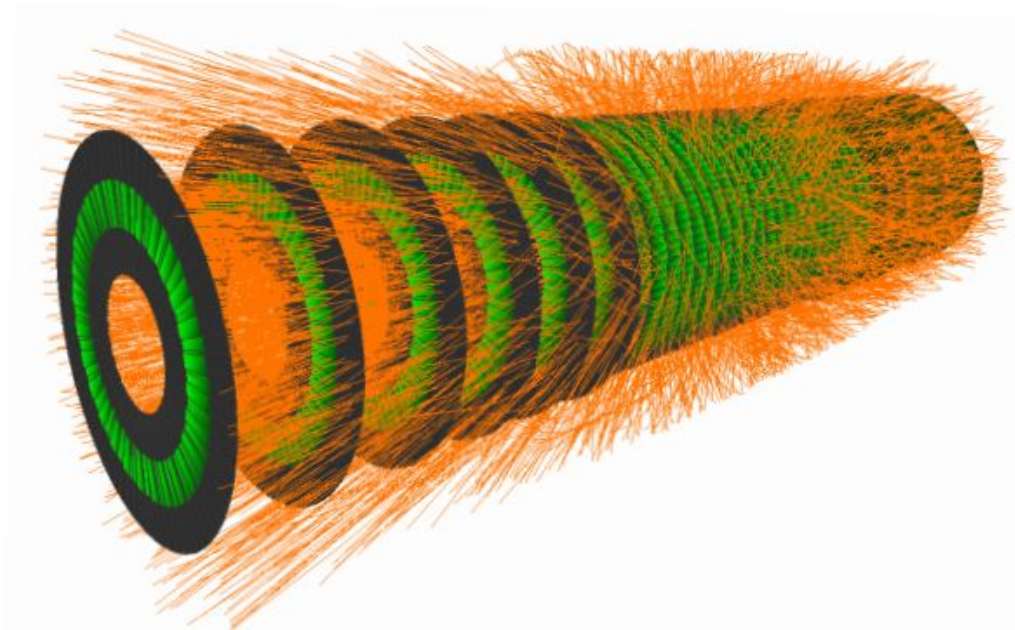


# DNA Sequencing

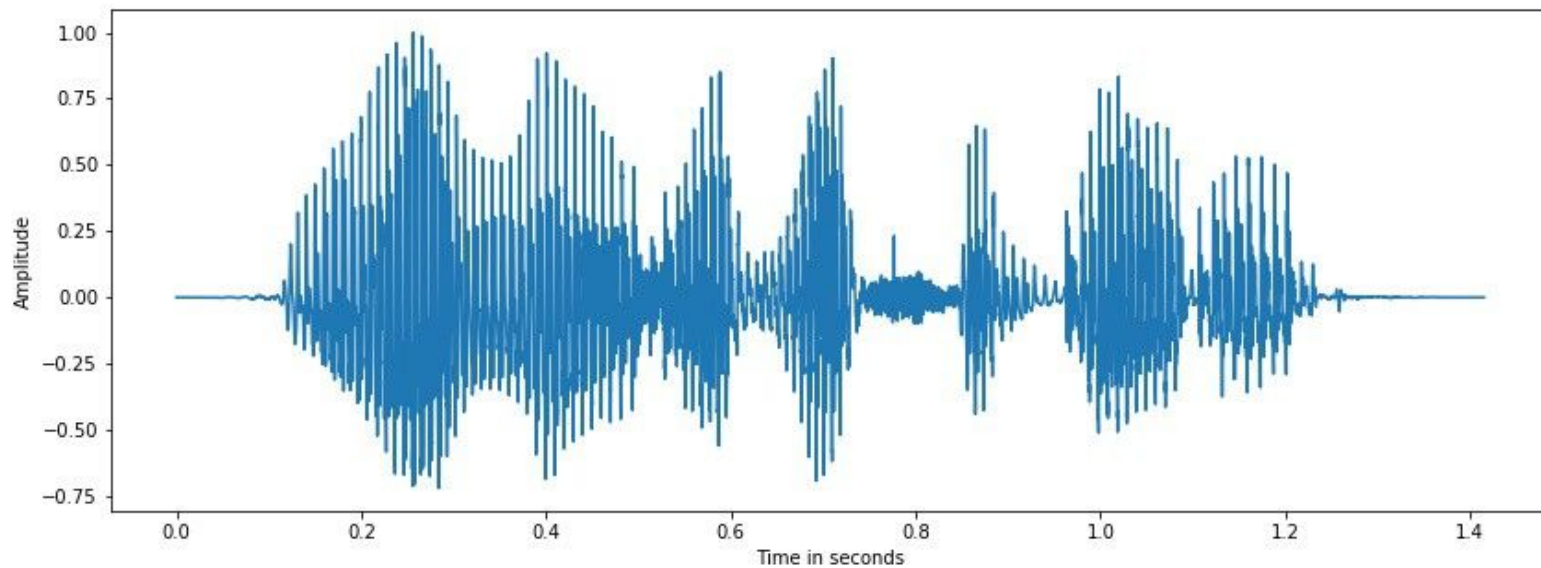


DNA sequence data from an automated sequencing machine

# Particle Track Reconstruction



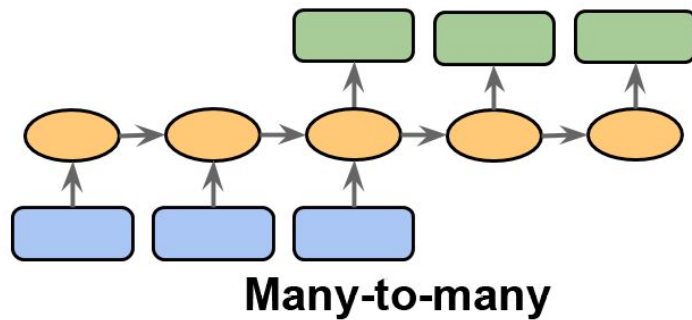
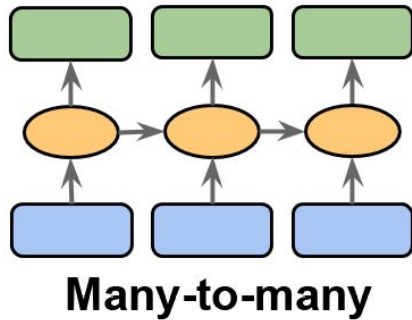
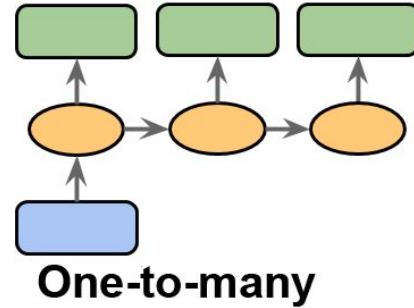
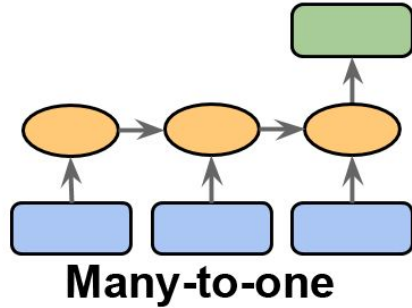
# Speech Analysis



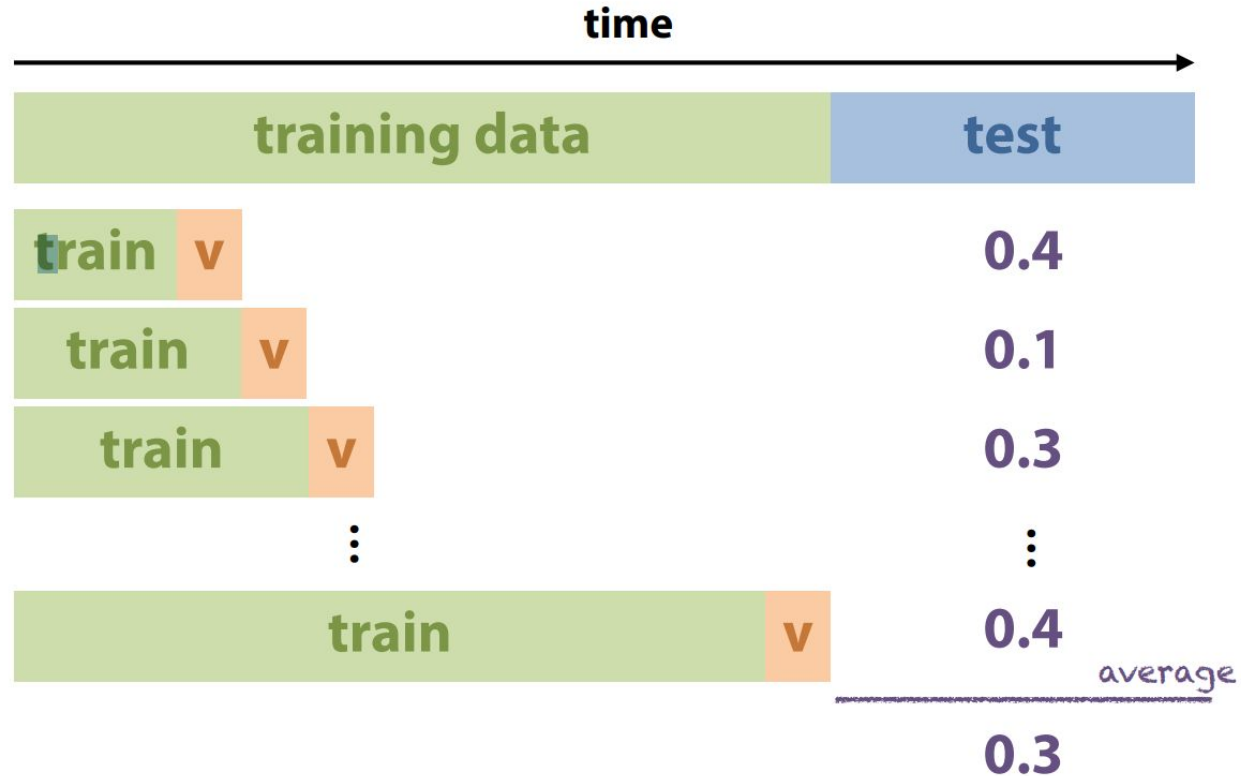
Try this sketching [app](#)



# Type of Sequences



# Learning from Sequences



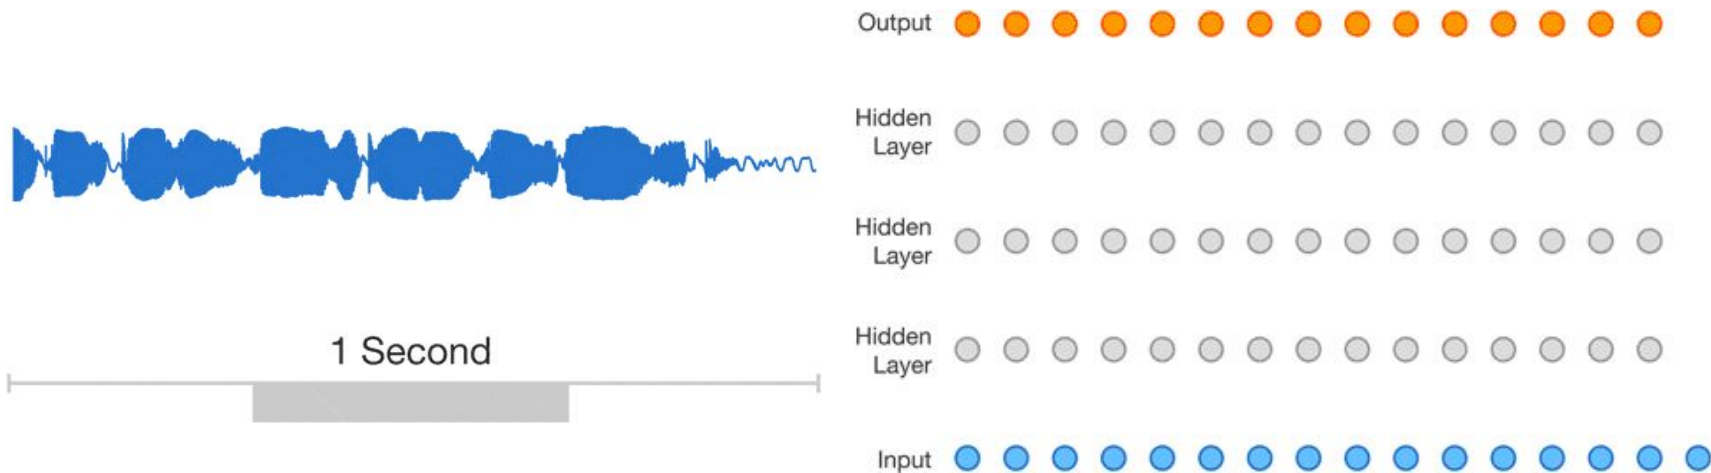
# Why not using a CNN?

- Extensions of 1D CNN.

Example: time series of  $N$  steps, each step with  $c$  features.

CNNs can be extended easily to other domains having grid-like structure of various dimensions. For example, consider a time-series of  $n$  steps, each step having  $c$  features (e.g.,  $c$  different readings from different sensors). We would need masked convolution (dilated)

# Example: WaveNet

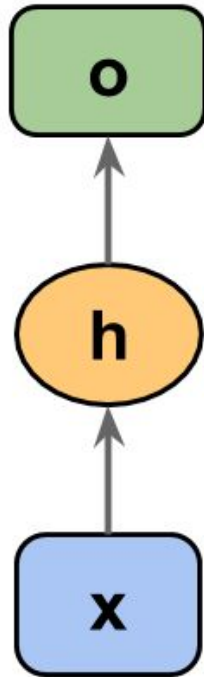


[WaveNet: A Generative Model for Raw Audio](#)

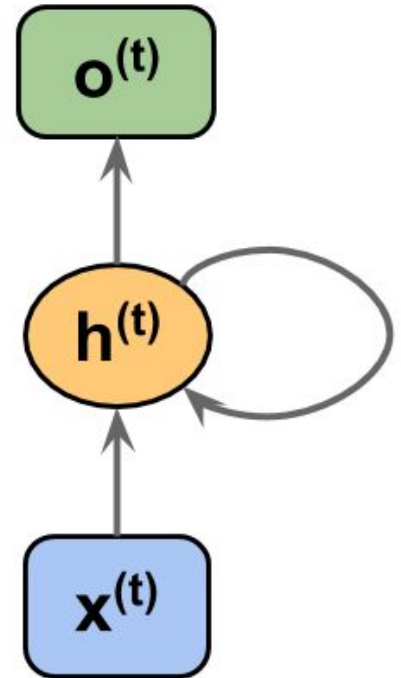
[DeepMind Blog Post](#)

# Recurrent Neural Networks

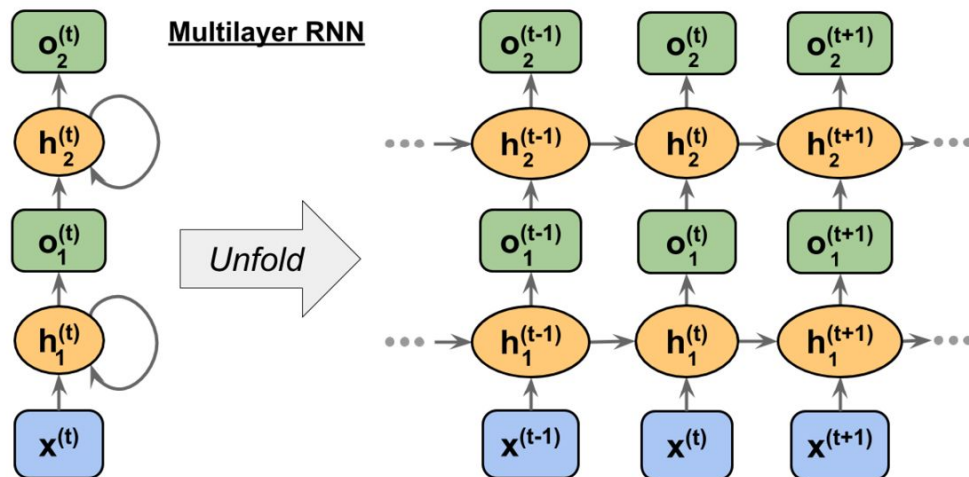
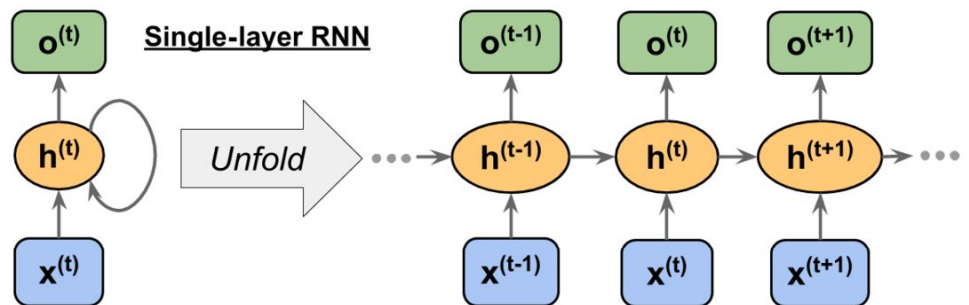
A standard  
feedforward  
network



Recurrent  
neural  
network



# Unfolding the RNN



# Elman Network (1990)

1. Start with hidden state:

$$h_0 = 0$$

2. Update with new state

$$h_t = \text{ReLU}(\mathbf{w}_{xh}x_t + \mathbf{w}_{hh}h_{t-1} + b_h)$$

3. Final prediction

$$y_T = \mathbf{w}_{hy}h_T + b_y.$$

# Elman Network (1990)

1. Start with hidden state:

$$h_0 = 0$$

2. Update with new state

$$h_t = \text{ReLU}(\mathbf{w}_{xh}x_t + \mathbf{w}_{hh}h_{t-1} + b_h)$$

3. Final prediction

$$y_T = \mathbf{w}_{hy}h_T + b_y.$$

```
class ElmanNet(nn.Module):  
  
    def __init__(self, size_input, size_hidden, size_output):  
  
        super(ElmanNet, self).__init__()  
  
        self.fc_x2h = nn.Linear(size_input, size_hidden)  
  
        self.fc_h2h = nn.Linear(size_hidden, size_hidden, bias=False)  
  
        self.fc_h2y = nn.Linear(size_recurrent, size_output)  
  
    def forward(self, x):  
  
        h = x.new_zeros(1, self.fc_h2y.weight.size(1))  
  
        for t in range(x.size(0)):  
  
            h = torch.relu(self.fc_x2h(x[t,:]) + self.fc_h2h(h))  
  
        return self.fc_h2y(h)
```

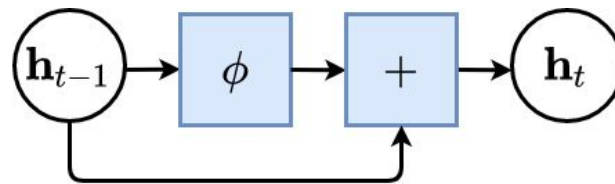


# Elman Network in PyTorch

```
class ElmanNet(nn.Module):  
  
    def __init__(self, size_input, size_hidden, size_output):  
        super(ElmanNet, self).__init__()  
  
        self.fc_x2h = nn.Linear(size_input, size_hidden)  
  
        self.fc_h2h = nn.Linear(size_hidden, size_hidden, bias=False)  
  
        self.fc_h2y = nn.Linear(size_recurrent, size_output)  
  
    def forward(self, x):  
  
        h = x.new_zeros(1, self.fc_h2y.weight.size(1))  
  
        for t in range(x.size(0)):  
            h = torch.relu(self.fc_x2h(x[t,:]) + self.fc_h2h(h))  
  
        return self.fc_h2y(h)
```

```
rnn = ElmanNet(size_input=10, size_hidden=50, size_output=2)  
  
cross_entropy = nn.CrossEntropyLoss()  
  
optimizer = torch.optim.Adam(rnn.parameters(), lr=learning_rate)  
  
for k in range(data.size()):  
    x,label = data.get_batch()  
  
    y = rnn(x)  
  
    loss = cross_entropy(y, label)  
  
    optimizer.zero_grad()  
  
    loss.backward()  
  
    optimizer.step()
```

# Gating



- Gates decide which information to go through.  
They are composed out of a sigmoid  $\sigma$  neural net layer and a pointwise multiplication operation.
- The sigmoid outputs numbers between zero and one, zero means “let nothing through,” while a value of one means “let everything through!”

# Gating Implementation

Update hidden state proposal: (same as Elman)

$$\bar{h}_t = \text{ReLU}(\mathbf{w}_{xh}x_t + \mathbf{w}_{hh}h_{t-1} + b_h)$$

Forget gate:

$$z_t = \sigma(\mathbf{w}_{xz}x_t + \mathbf{w}_{hz}h_{t-1} + b_z)$$

Hidden State:

$$h_t = z_t \odot h_{t-1} + (1 - z_t) \odot \bar{h}_t$$

# Gating Implementation

Update hidden state proposal: (same as Elman)

$$\bar{h}_t = \text{ReLU}(\mathbf{w}_{xh}x_t + \mathbf{w}_{hh}h_{t-1} + b_h)$$

Forget gate:

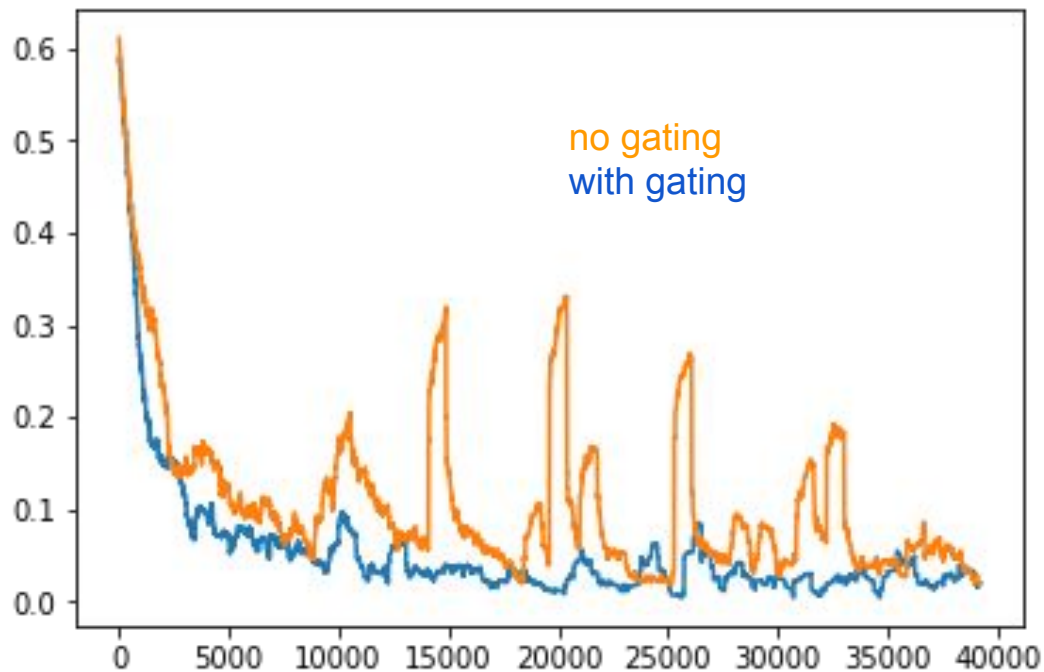
$$z_t = \sigma(\mathbf{w}_{xz}x_t + \mathbf{w}_{hz}h_{t-1} + b_z)$$

Hidden State:

$$h_t = z_t \odot h_{t-1} + (1 - z_t) \odot \bar{h}_t$$

```
class ElmanNetGating(nn.Module):  
  
    def __init__(self, size_input, size_hidden, size_output):  
  
        super(ElmanNetGating, self).__init__()  
  
        self.fc_x2h = nn.Linear(size_input, size_hidden)  
  
        self.fc_h2h = nn.Linear(size_hidden, size_hidden, bias=False)  
  
        self.fc_x2z = nn.Linear(size_input, size_hidden)  
  
        self.fc_h2z = nn.Linear(size_hidden, size_hidden, bias=False)  
  
        self.fc_h2y = nn.Linear(size_hidden, size_output)  
  
    def forward(self, x):  
  
        h = x.new_zeros(1, self.fc_h2y.weight.size(1))  
  
        for t in range(x.size(0)):  
  
            z = torch.sigmoid(self.fc_x2z(x[t,:]) + self.fc_h2z(h))  
  
            hb = torch.relu(self.fc_x2h(x[t,:]) + self.fc_h2h(h))  
  
            h = z * h + (1-z) * hb  
  
        return self.fc_h2y(h)
```

# Training with Gating



Loss curve

# Gated Recurrent Units (GRU)

Update hidden state proposal:

$$\bar{h}_t = \tanh(\mathbf{w}_{xh}x_t + \mathbf{w}_{hh}(r_t \odot h_{t-1}) + b_h)$$

Forget gate:

$$z_t = \sigma(\mathbf{w}_{xz}x_t + \mathbf{w}_{hz}h_{t-1} + b_z)$$

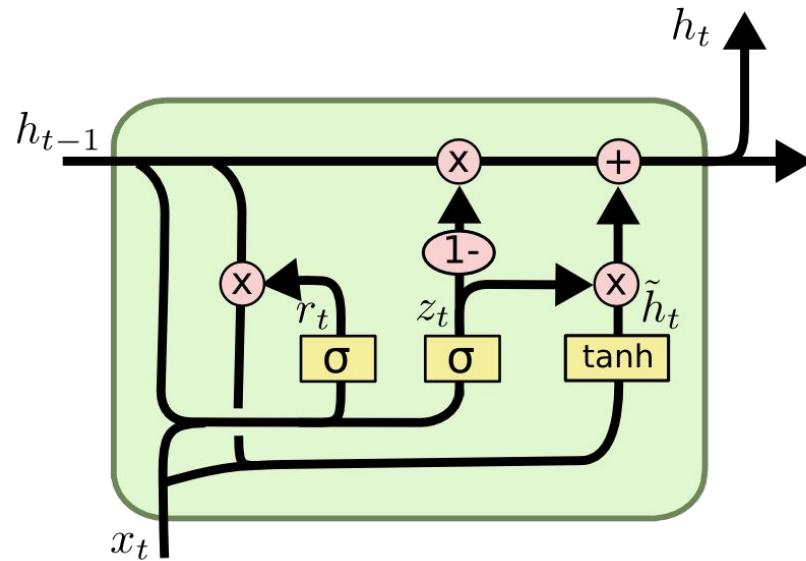
Reset gate:

$$r_t = \sigma(\mathbf{w}_{xr}x_t + \mathbf{w}_{hr}h_{t-1} + b_r)$$

Hidden State:

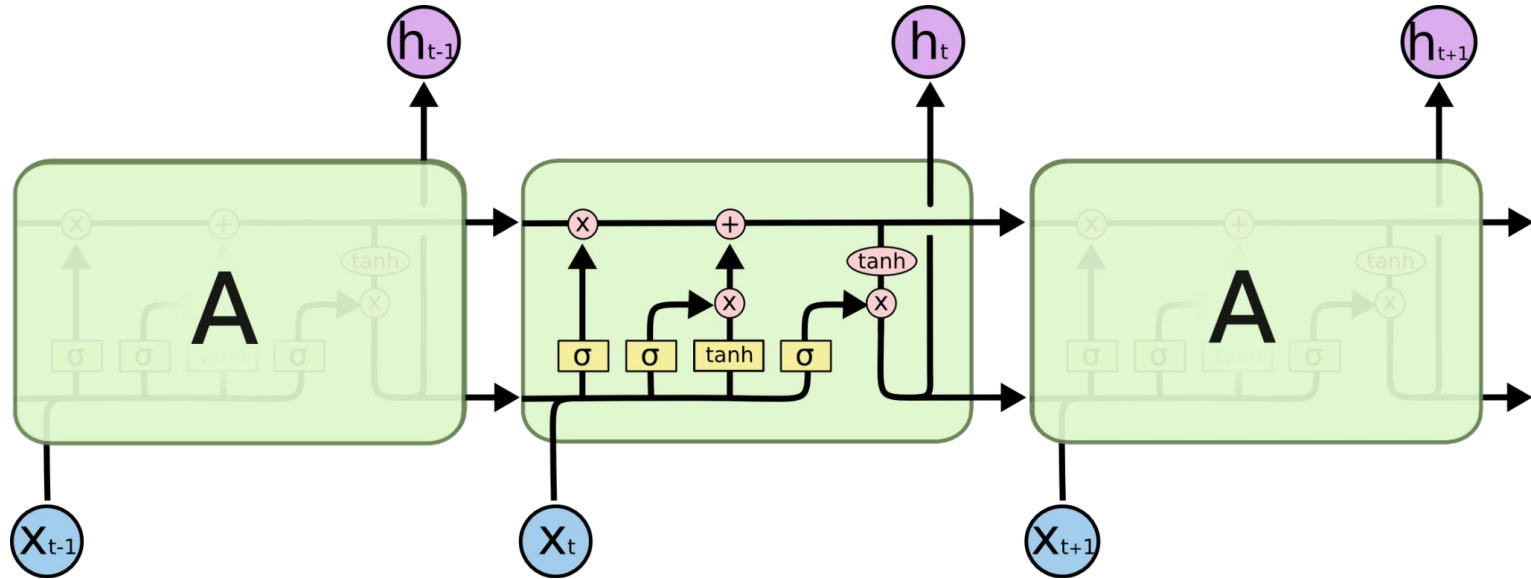
$$h_t = z_t \odot h_{t-1} + (1 - z_t) \odot \bar{h}_t$$

# GRU



[Cho et. al \(2014\)](#)

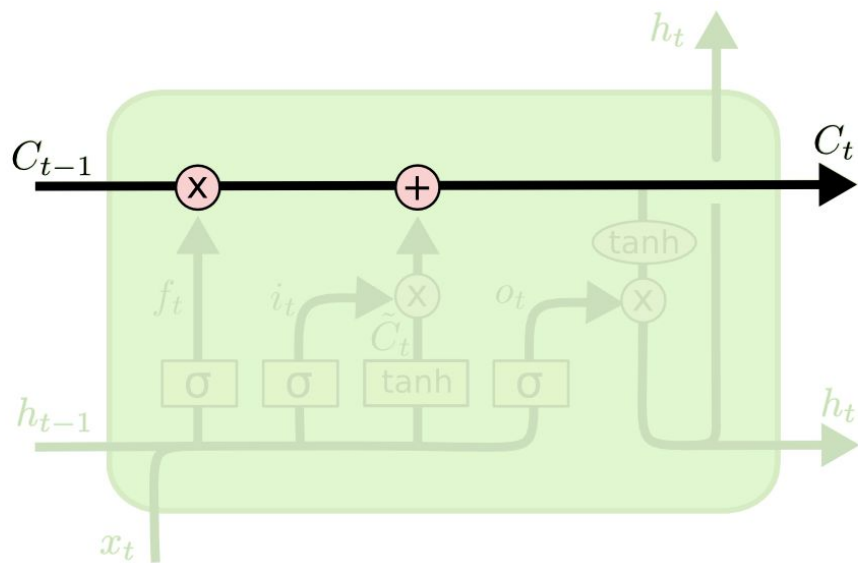
# Long Short-Term Memory (LSTM)



[Hochreiter & Schmidhuber \(1997\)](#)

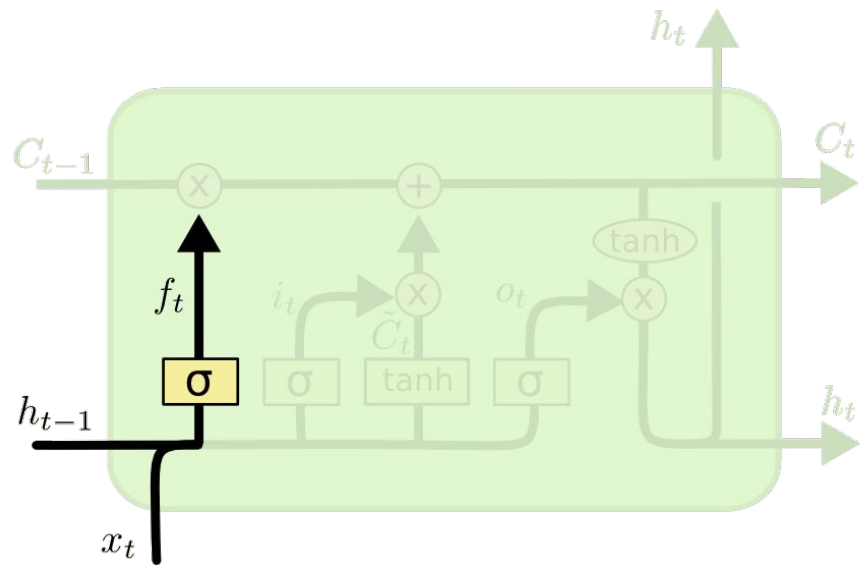


# LSTM



Cell State

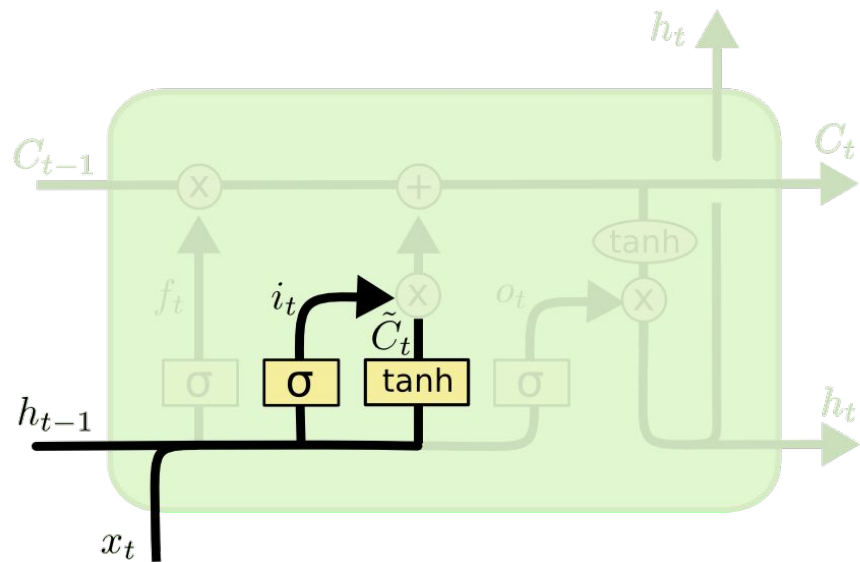
# LSTM



Forget gate layer

$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$

# LSTM

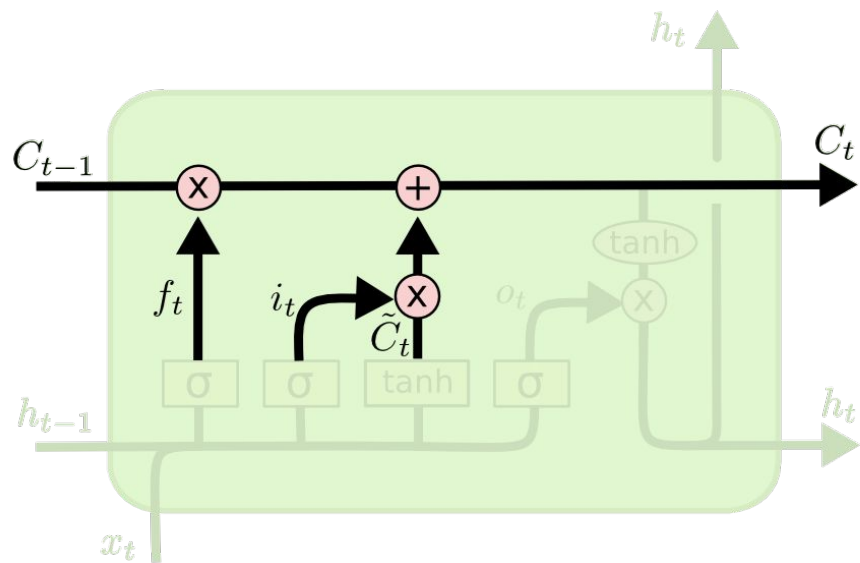


Input Gate Layer

$$i_t = \sigma (W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

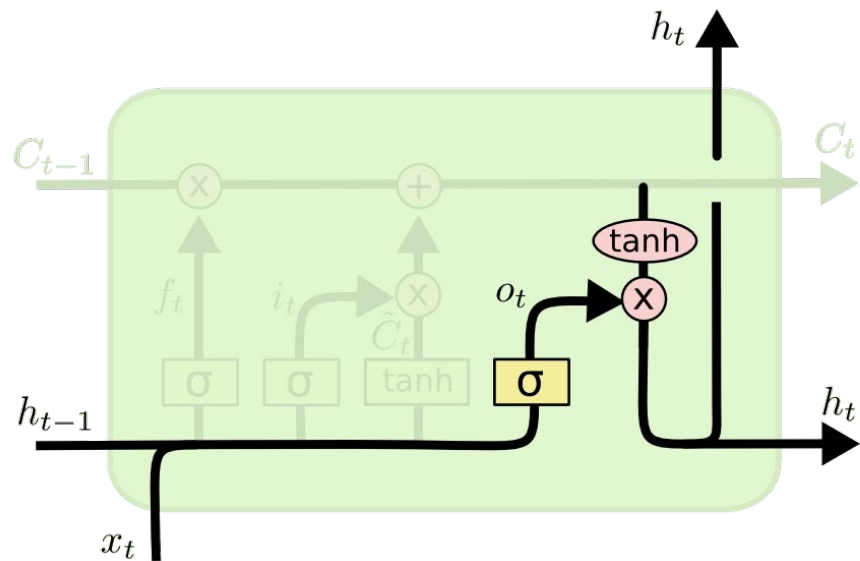
# LSTM



Update Cell State

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

# LSTM



Output gate

$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

# LSTM in PyTorch

```
class MyLSTM(nn.Module):

    def __init__(self, size_input, size_hidden, num_layers, size_output):

        super(MyLSTM, self).__init__()

        self.lstm = nn.LSTM(input_size=size_input, hidden_size=size_hidden, num_layers=num_layers)

        self.fc_o2y = nn.Linear(size_hidden, size_output)

    def forward(self, x):

        x = x.unsqueeze(1) # expect a batch size (here is 1)

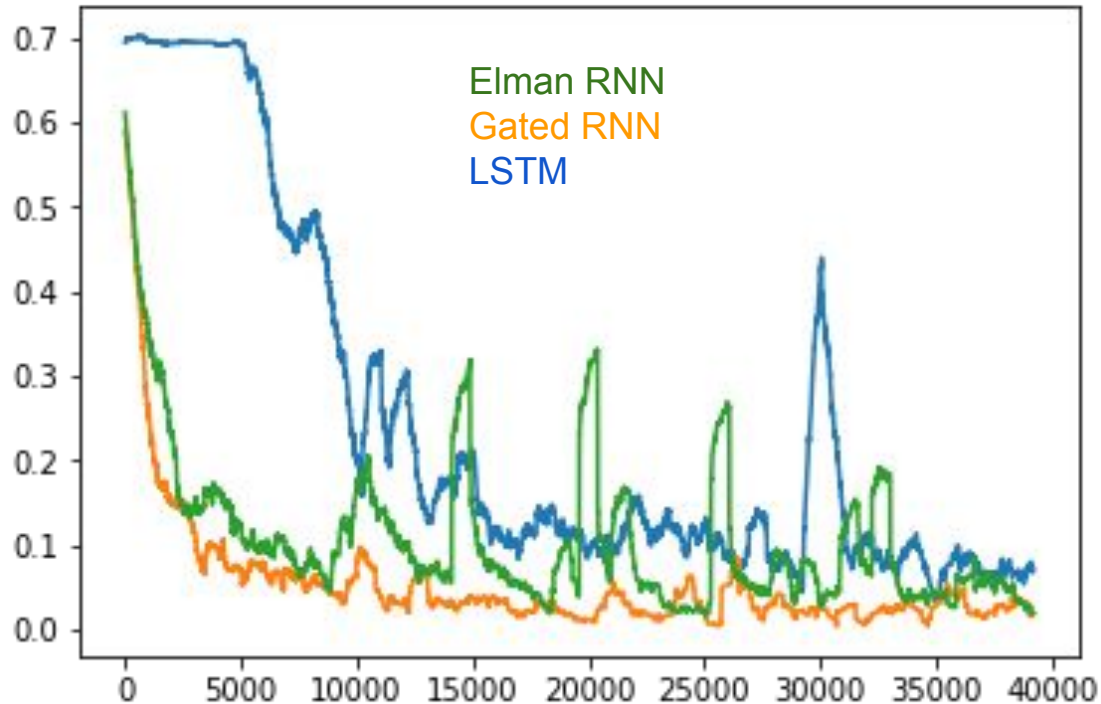
        output, _ = self.lstm(x)

        output = output.squeeze(1) # only last layer, shape (seq. len., bs, dim_recurrent) and drop the batch index

        output = output.narrow(0, output.size(0)-1,1) # keep only the last hidden variable

        return self.fc_o2y(F.relu(output)) # shape (1, dim_recurrent)
```

# Training with LSTM



Loss curve

# Resources

- Chris Olah: [Understanding LSTM Networks](#)
- Andrej Karpathy: [The Unreasonable Effectiveness of Recurrent Neural Networks](#)