ML: Some Guidelines

University of Victoria - PHYS-555

1 - Problem Formulation

- What am I trying to solve?
    - do I need ML? supervised? regression? would classify those objects really give me a better understanding of the physics?
- Which information do I have access to answer this problem?
    - structured data? noisy? metadata? is there hidden information?
    - how can I inform my model with induction biases that I perceive on the data?

2 - Become one with the data

- **Estimate**:
  - How much data do I have access to? How hard is it to get more?
  - What is the data structure: sequences, images, tables. Any known relations?
  - What was the selection process for the data?
  - Can I possibly do data augmentation | simulation?
- **Explore**: plot distributions, correlations, missing, imbalance, get an intuition!
- **Label**: what kind of labels I have access? Are they noisy?
- **Preprocess**: will I need rescaling, do I need feature engineering?

3 - Modelling and interpretation

- **Algorithm**:
  - select algorithms according to the data (tabular, sequential, images, 3D, irregular graphs,...)
  - start with a reliable baseline, i.e. RandomForest, ResNet18, …
  - be ready to pipeline several algorithms
- **Evaluation**: make sure you have metrics
- **Interpretation**: can I trace the results back to the data and act?
- **Robustness:** if I share my pre-trained model, how would it fail?
- **Hyper-parameters**: the gain is often small, reserve for last.

Time Management

Typical perception

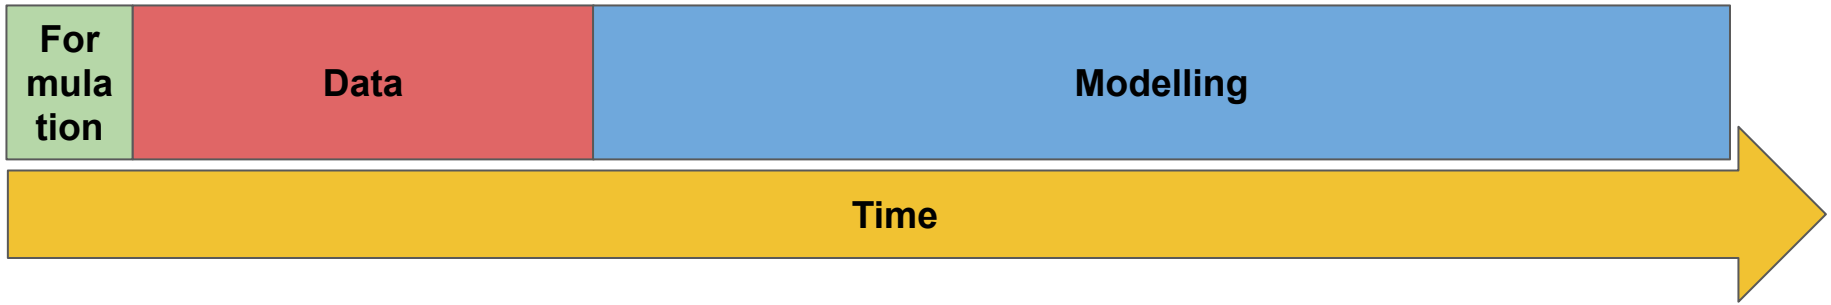

Time Management

Perceived impact years after results

Be ready to seriously revisit your initial time estimates